



# Hierarchical Clustering

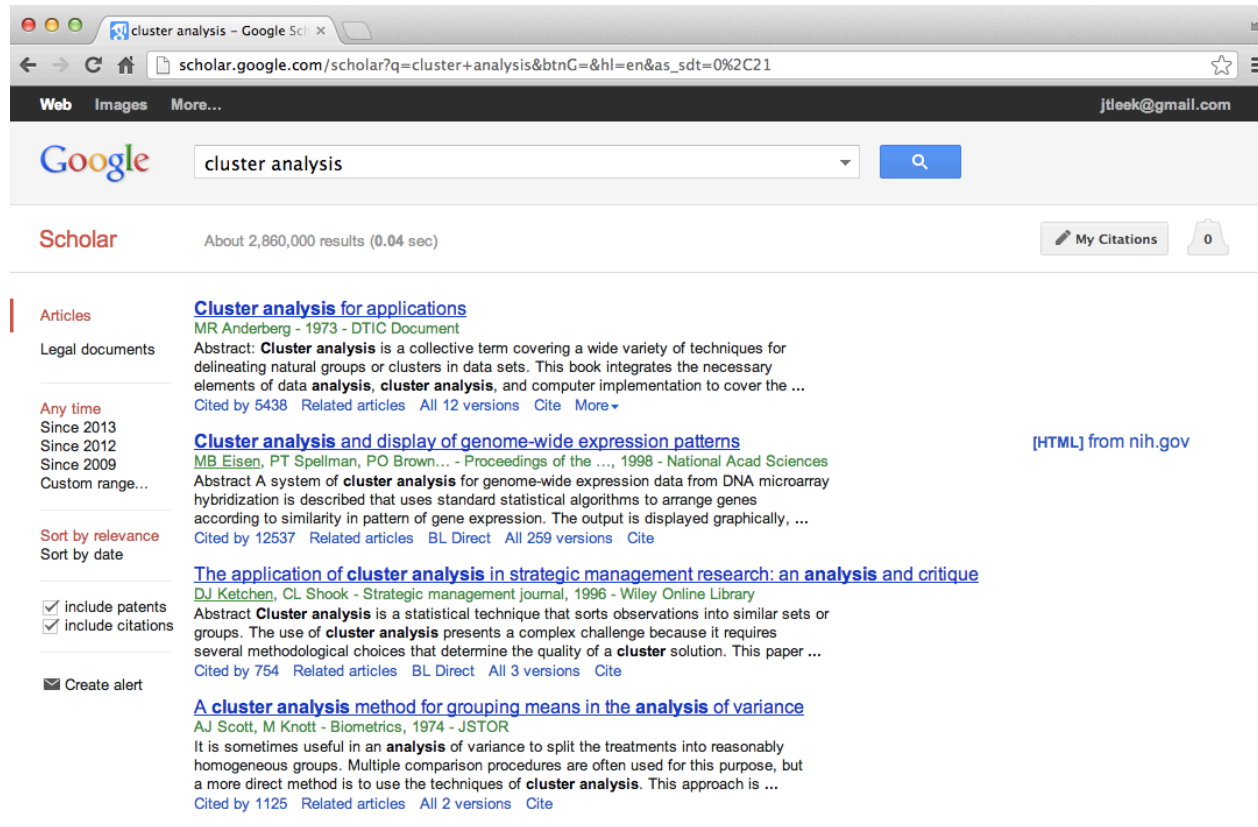
Roger D. Peng, Associate Professor of Biostatistics  
Johns Hopkins Bloomberg School of Public Health

# Can we find things that are close together?

Clustering organizes things that are **close** into groups

- How do we define close?
- How do we group things?
- How do we visualize the grouping?
- How do we interpret the grouping?

# Hugely important/impactful



[http://scholar.google.com/scholar?hl=en&q=cluster+analysis&btnG=&as\\_sdt=1%2C21&as\\_sdt=1%2C21](http://scholar.google.com/scholar?hl=en&q=cluster+analysis&btnG=&as_sdt=1%2C21&as_sdt=1%2C21)

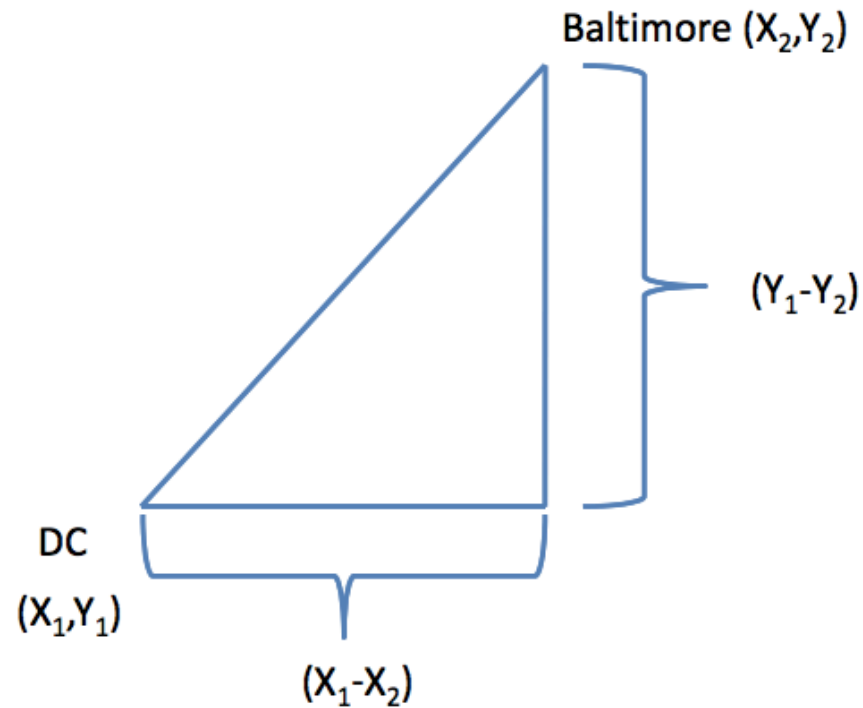
# Hierarchical clustering

- An agglomerative approach
  - Find closest two things
  - Put them together
  - Find next closest
- Requires
  - A defined distance
  - A merging approach
- Produces
  - A tree showing how close things are to each other

# How do we define close?

- Most important step
  - Garbage in -> garbage out
- Distance or similarity
  - Continuous - euclidean distance
  - Continuous - correlation similarity
  - Binary - manhattan distance
- Pick a distance/similarity that makes sense for your problem

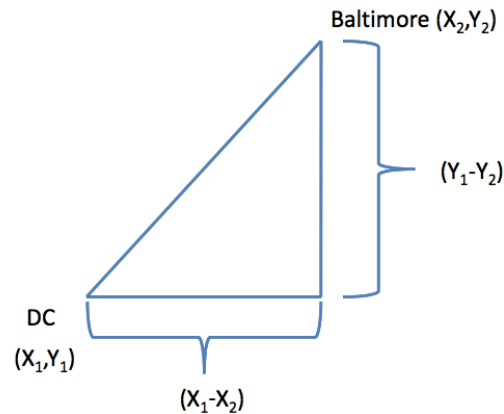
# Example distances - Euclidean



<http://rafalab.jhsph.edu/688/lec/lecture5-clustering.pdf>

# Example distances - Euclidean

$$\sqrt{(X_1 - X_2)^2 + (Y_1 - Y_2)^2}$$

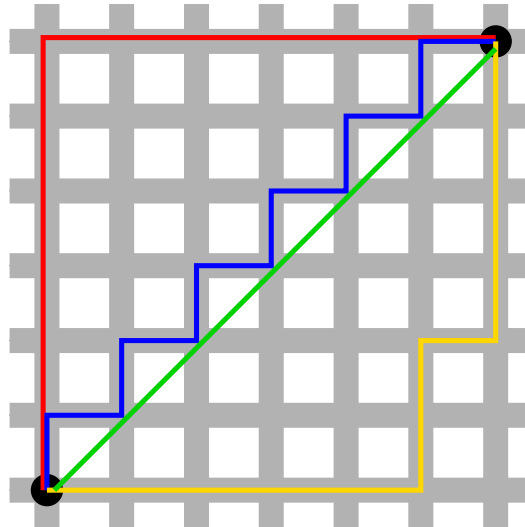


In general:

$$\sqrt{(A_1 - A_2)^2 + (B_1 - B_2)^2 + \dots + (Z_1 - Z_2)^2}$$

<http://rafalab.jhsph.edu/688/lec/lecture5-clustering.pdf>

# Example distances - Manhattan



In general:

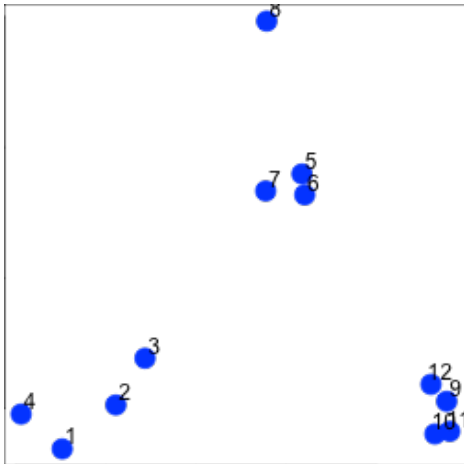
$$|A_1 - A_2| + |B_1 - B_2| + \dots + |Z_1 - Z_2|$$

[http://en.wikipedia.org/wiki/Taxicab\\_geometry](http://en.wikipedia.org/wiki/Taxicab_geometry)



# Hierarchical clustering - example

```
set.seed(1234)
par(mar = c(0, 0, 0, 0))
x <- rnorm(12, mean = rep(1:3, each = 4), sd = 0.2)
y <- rnorm(12, mean = rep(c(1, 2, 1), each = 4), sd = 0.2)
plot(x, y, col = "blue", pch = 19, cex = 2)
text(x + 0.05, y + 0.05, labels = as.character(1:12))
```



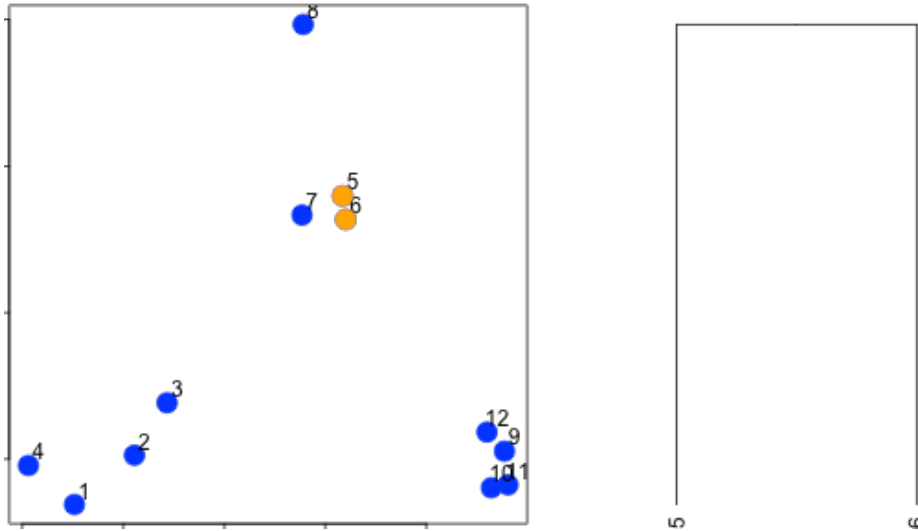
# Hierarchical clustering - dist

- Important parameters: *x, method*

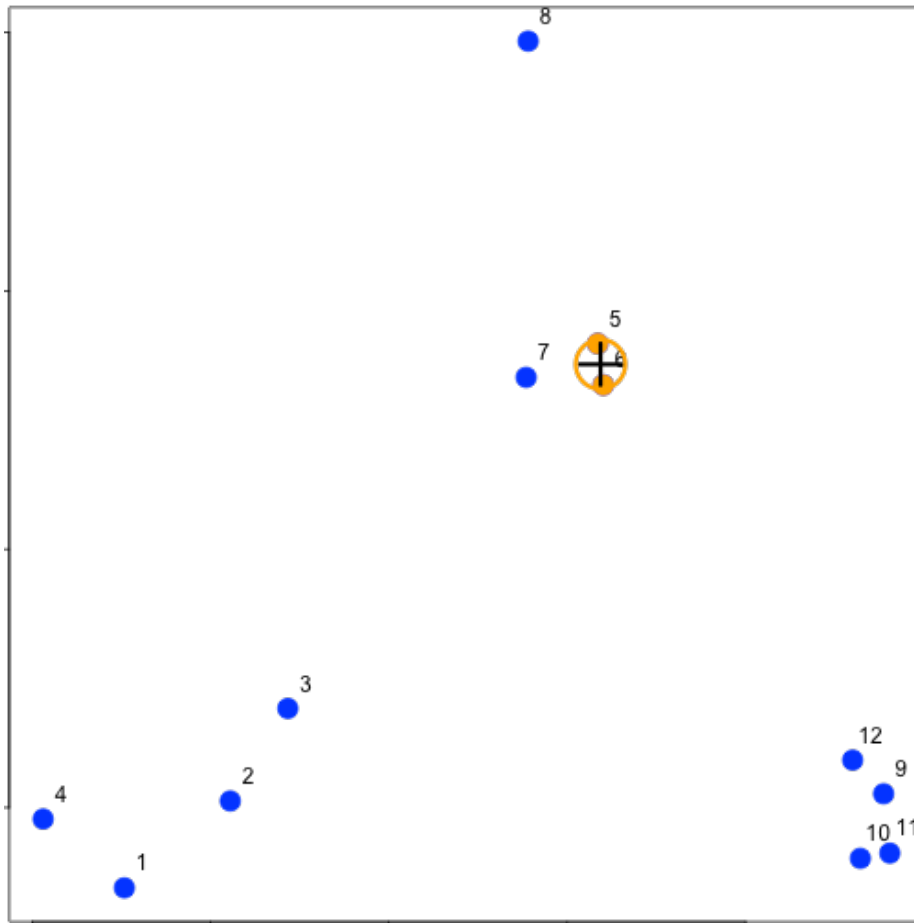
```
dataFrame <- data.frame(x = x, y = y)
dist(dataFrame)
```

```
##           1           2           3           4           5           6           7           8           9
## 2  0.34121
## 3  0.57494 0.24103
## 4  0.26382 0.52579 0.71862
## 5  1.69425 1.35818 1.11953 1.80667
## 6  1.65813 1.31960 1.08339 1.78081 0.08150
## 7  1.49823 1.16621 0.92569 1.60132 0.21110 0.21667
## 8  1.99149 1.69093 1.45649 2.02849 0.61704 0.69792 0.65063
## 9  2.13630 1.83168 1.67836 2.35676 1.18350 1.11500 1.28583 1.76461
## 10 2.06420 1.76999 1.63110 2.29239 1.23848 1.16550 1.32063 1.83518 0.14090
## 11 2.14702 1.85183 1.71074 2.37462 1.28154 1.21077 1.37370 1.86999 0.11624
## 12 2.05664 1.74663 1.58659 2.27232 1.07701 1.00777 1.17740 1.66224 0.10849
##           10           11
## 2
## 3
## 4
## 5
```

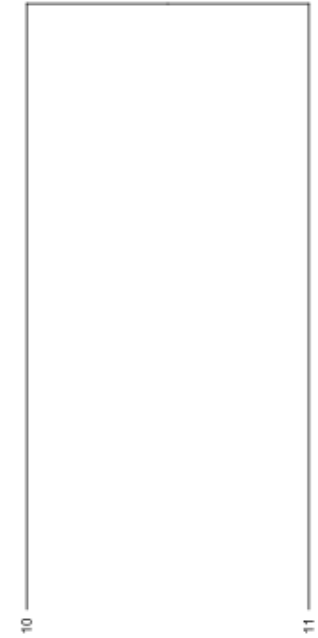
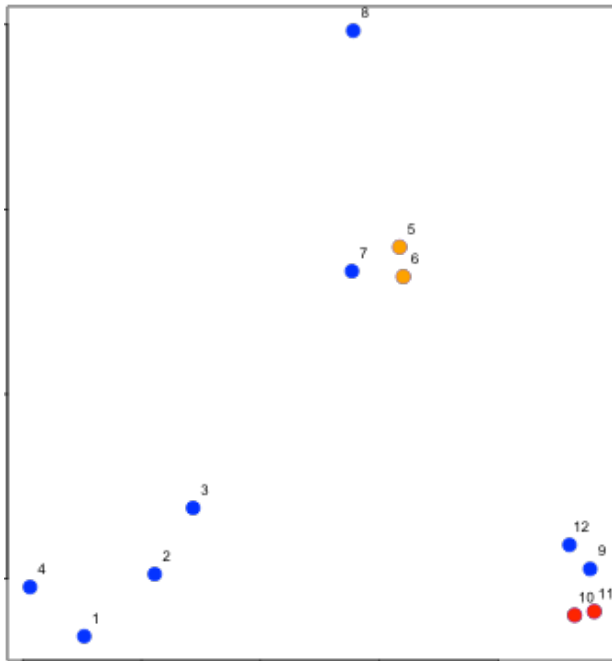
# Hierarchical clustering - #1



# Hierarchical clustering - #2

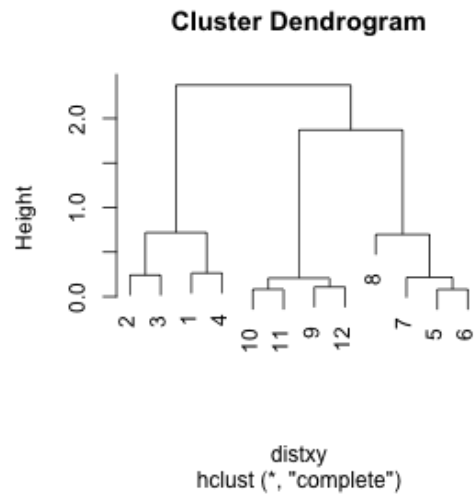


# Hierarchical clustering - #3



# Hierarchical clustering - hclust

```
dataFrame <- data.frame(x = x, y = y)
distxy <- dist(dataFrame)
hClustering <- hclust(distxy)
plot(hClustering)
```

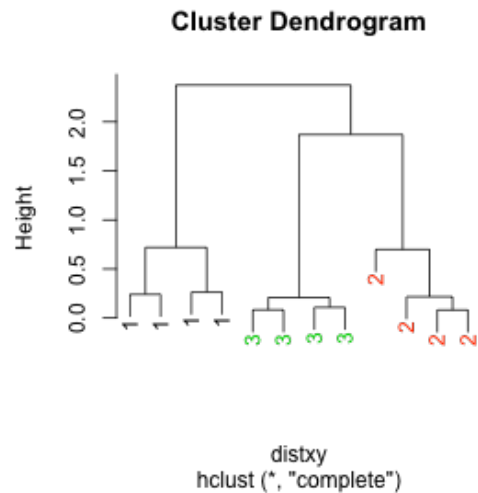


# Prettier dendrograms

```
myplclust <- function(hclust, lab = hclust$labels, lab.col = rep(1, length(hclust$labels)),
  hang = 0.1, ...) {
  ## modification of plclust for plotting hclust objects *in colour*! Copyright
  ## Eva KF Chan 2009 Arguments: hclust: hclust object lab: a character vector
  ## of labels of the leaves of the tree lab.col: colour for the labels;
  ## NA=default device foreground colour hang: as in hclust & plclust Side
  ## effect: A display of hierarchical cluster with coloured leaf labels.
  y <- rep(hclust$height, 2)
  x <- as.numeric(hclust$merge)
  y <- y[which(x < 0)]
  x <- x[which(x < 0)]
  x <- abs(x)
  y <- y[order(x)]
  x <- x[order(x)]
  plot(hclust, labels = FALSE, hang = hang, ...)
  text(x = x, y = y[hclust$order] - (max(hclust$height) * hang), labels = lab[hclust$order],
    col = lab.col[hclust$order], srt = 90, adj = c(1, 0.5), xpd = NA, ...)
}
```

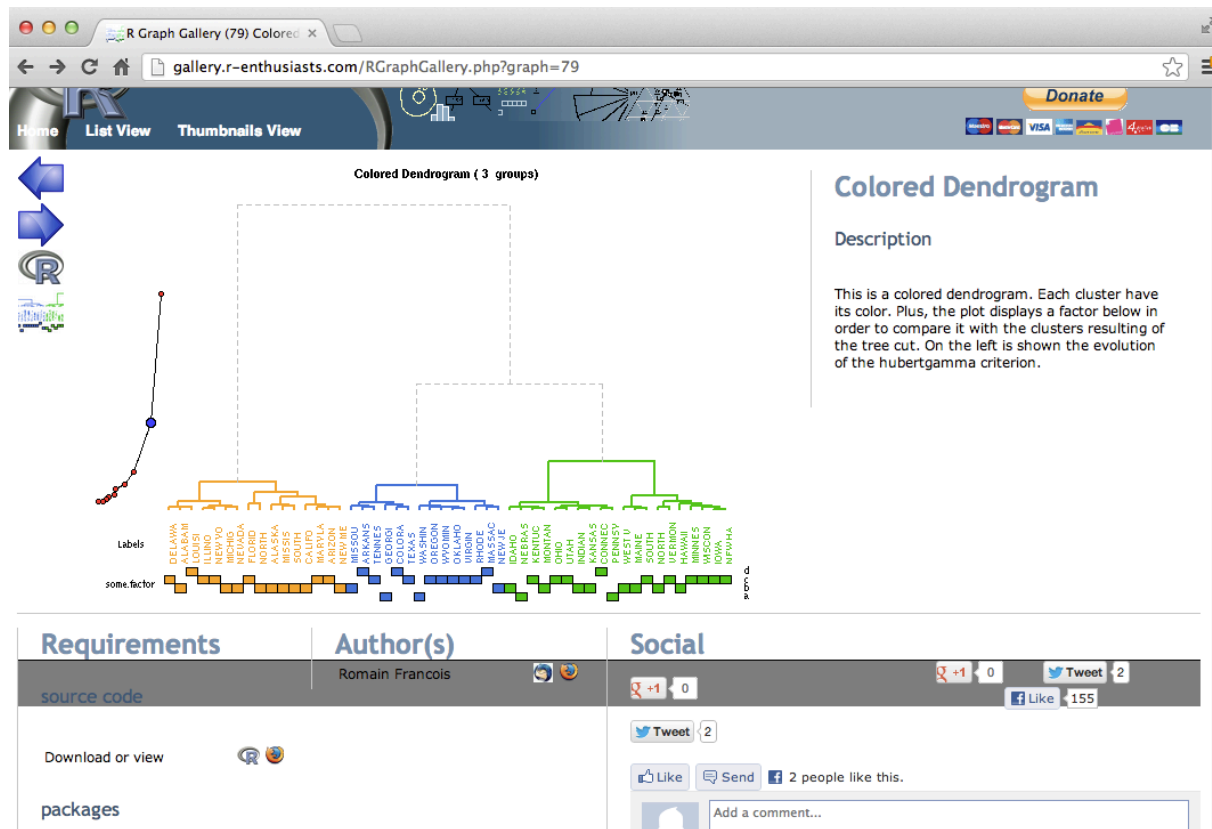
# Pretty dendrograms

```
dataFrame <- data.frame(x = x, y = y)
distxy <- dist(dataFrame)
hClustering <- hclust(distxy)
myplclust(hClustering, lab = rep(1:3, each = 4), lab.col = rep(1:3, each = 4))
```



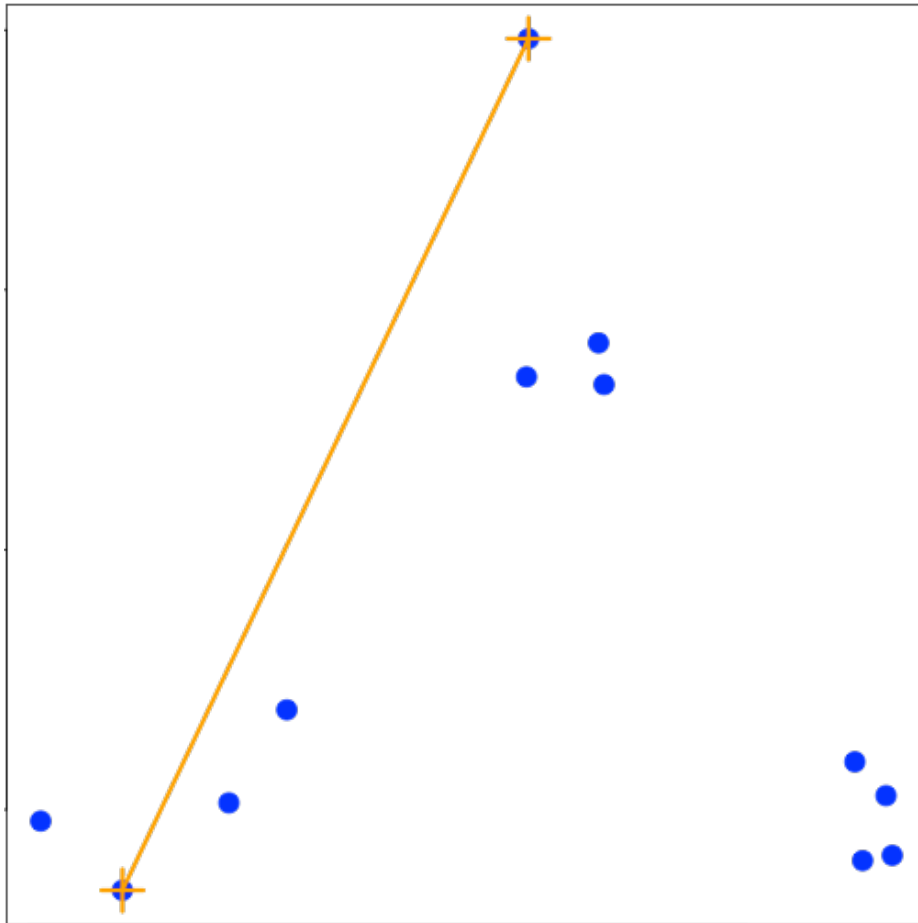


# Even Prettier dendrograms

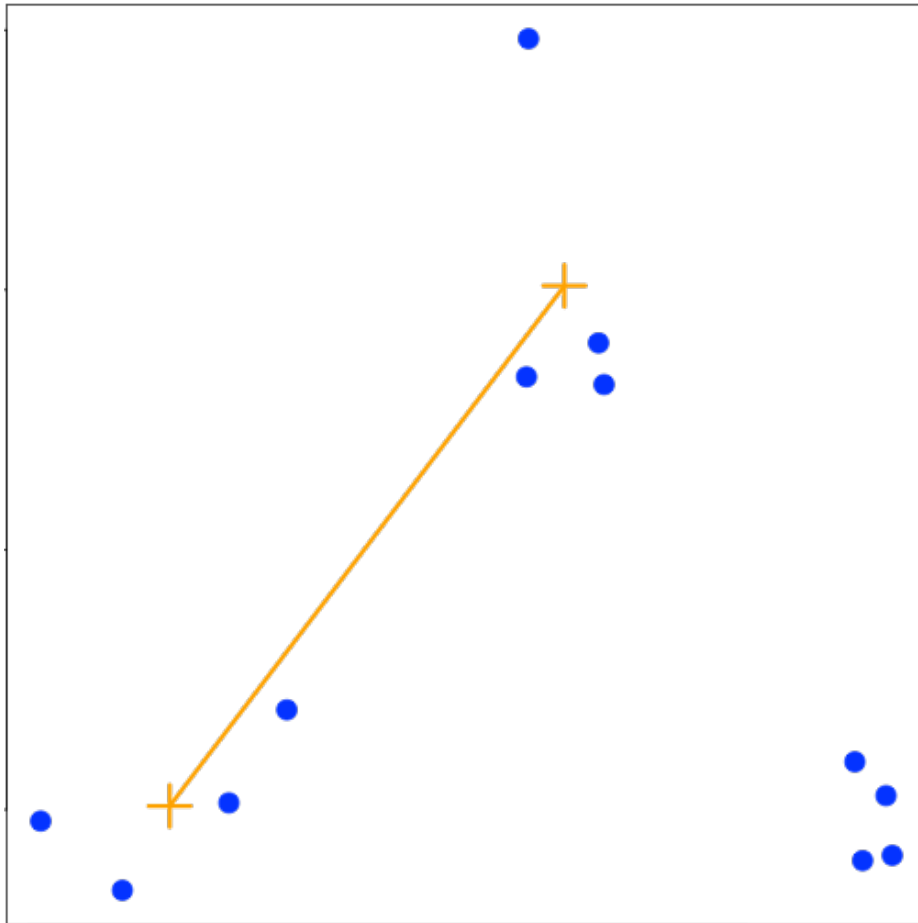


<http://gallery.r-enthusiasts.com/RGraphGallery.php?graph=79>

# Merging points - complete

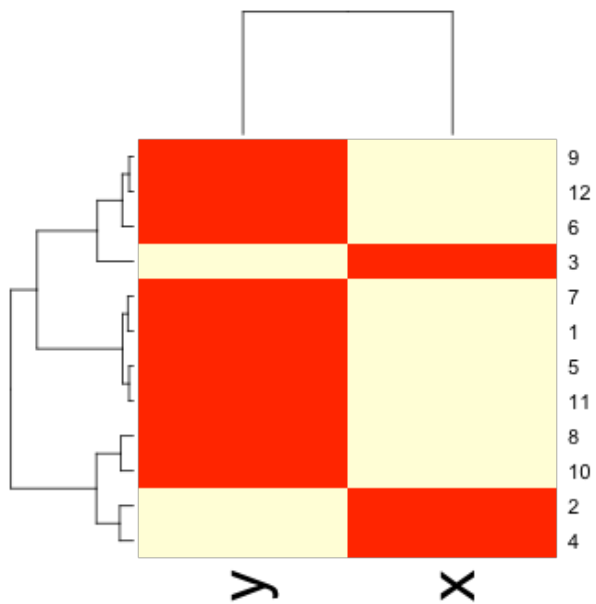


# Merging points - average



# heatmap()

```
dataFrame <- data.frame(x = x, y = y)
set.seed(143)
dataMatrix <- as.matrix(dataFrame)[sample(1:12), ]
heatmap(dataMatrix)
```



# Notes and further resources

- Gives an idea of the relationships between variables/observations
- The picture may be unstable
  - Change a few points
  - Have different missing values
  - Pick a different distance
  - Change the merging strategy
  - Change the scale of points for one variable
- But it is deterministic
- Choosing where to cut isn't always obvious
- Should be primarily used for exploration
- [Rafa's Distances and Clustering Video](#)
- [Elements of statistical learning](#)



# K-means Clustering

Roger D. Peng, Associate Professor of Biostatistics  
Johns Hopkins Bloomberg School of Public Health

# Can we find things that are close together?

- How do we define close?
- How do we group things?
- How do we visualize the grouping?
- How do we interpret the grouping?

# How do we define close?

- Most important step
  - Garbage in  $\longrightarrow$  garbage out
- Distance or similarity
  - Continuous - euclidean distance
  - Continuous - correlation similarity
  - Binary - manhattan distance
- Pick a distance/similarity that makes sense for your problem

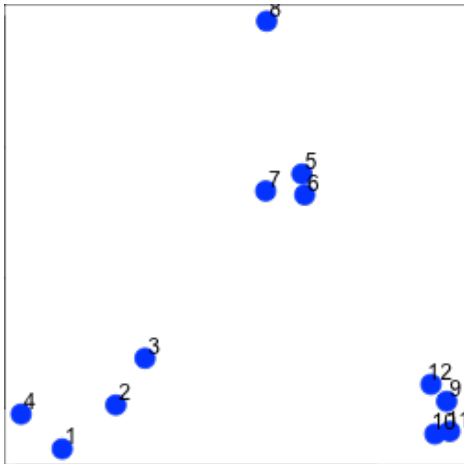


# K-means clustering

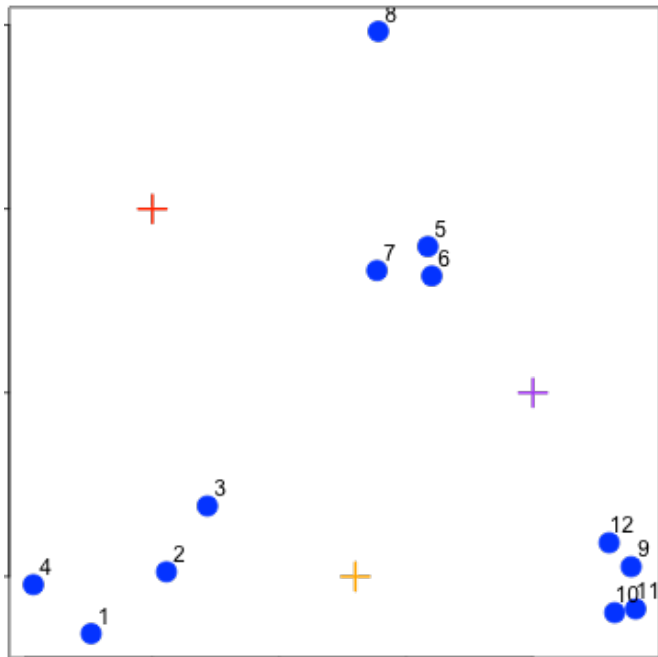
- A partitioning approach
  - Fix a number of clusters
  - Get "centroids" of each cluster
  - Assign things to closest centroid
  - Reclaculate centroids
- Requires
  - A defined distance metric
  - A number of clusters
  - An initial guess as to cluster centroids
- Produces
  - Final estimate of cluster centroids
  - An assignment of each point to clusters

# K-means clustering - example

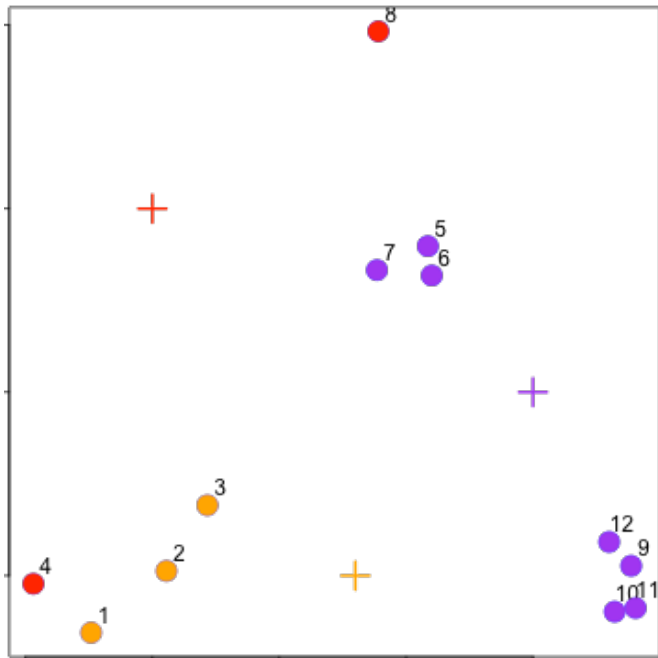
```
set.seed(1234)
par(mar = c(0, 0, 0, 0))
x <- rnorm(12, mean = rep(1:3, each = 4), sd = 0.2)
y <- rnorm(12, mean = rep(c(1, 2, 1), each = 4), sd = 0.2)
plot(x, y, col = "blue", pch = 19, cex = 2)
text(x + 0.05, y + 0.05, labels = as.character(1:12))
```



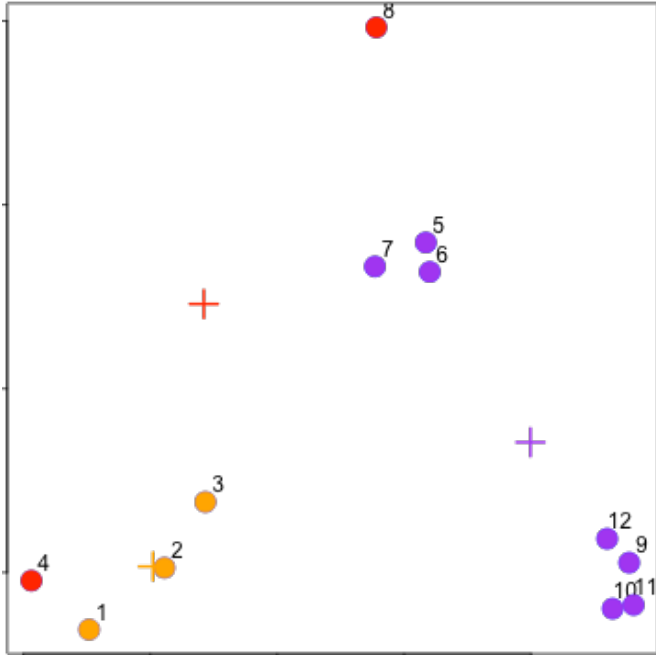
# K-means clustering - starting centroids



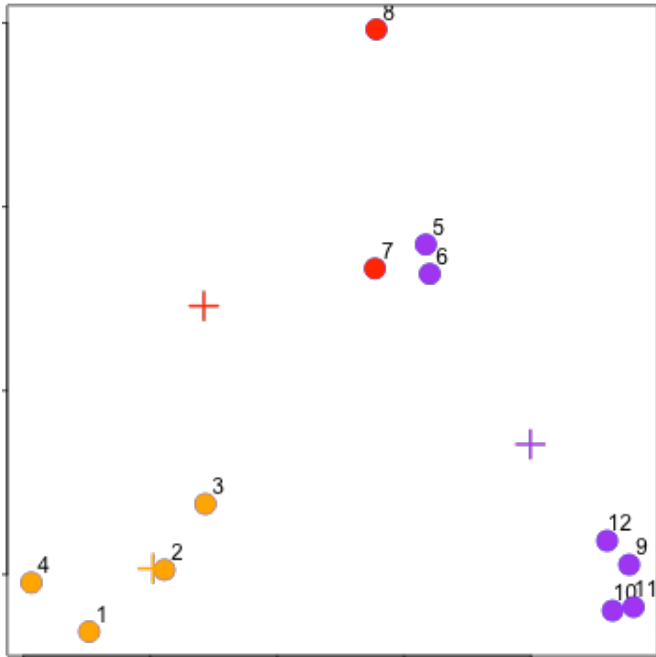
# K-means clustering - assign to closest centroid



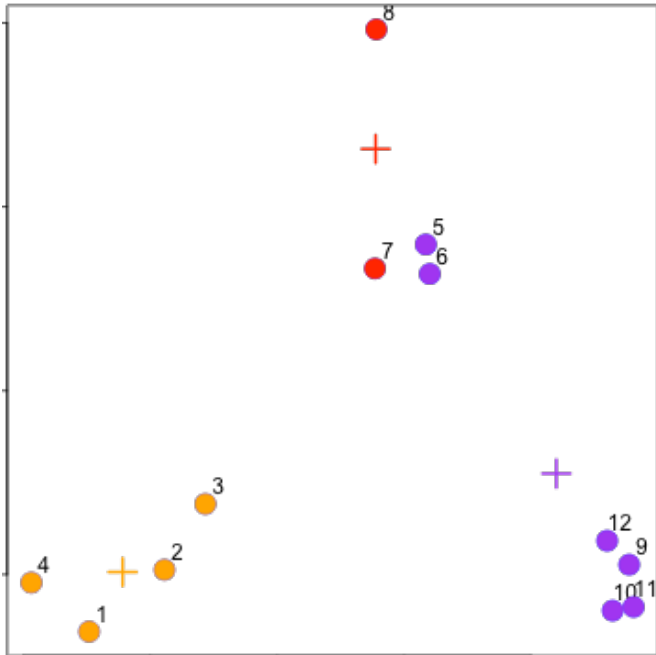
# K-means clustering - recalculate centroids



# K-means clustering - reassign values



# K-means clustering - update centroids



# kmeans ( )

- Important parameters: *x*, *centers*, *iter.max*, *nstart*

```
dataFrame <- data.frame(x, y)
kmeansObj <- kmeans(dataFrame, centers = 3)
names(kmeansObj)
```

```
## [1] "cluster"      "centers"      "totss"       "withinss"
## [5] "tot.withinss" "betweenss"    "size"        "iter"
## [9] "ifault"
```

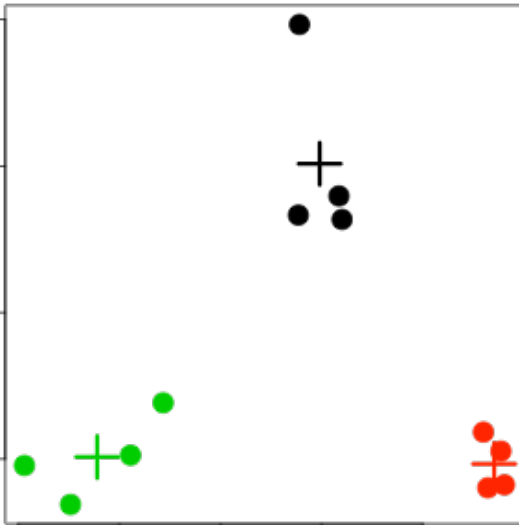
```
kmeansObj$cluster
```

```
## [1] 3 3 3 3 1 1 1 1 2 2 2 2
```



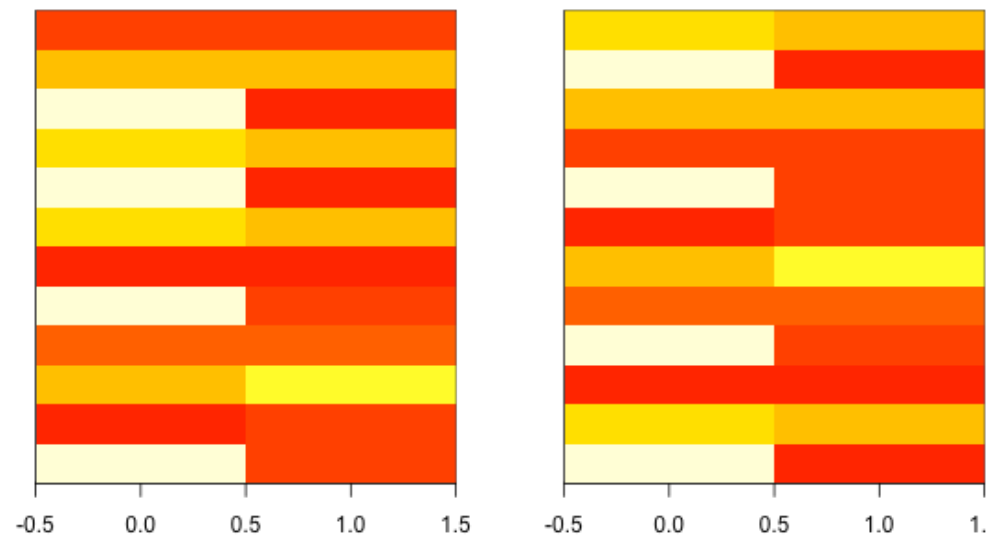
# kmeans ( )

```
par(mar = rep(0.2, 4))  
plot(x, y, col = kmeansObj$cluster, pch = 19, cex = 2)  
points(kmeansObj$centers, col = 1:3, pch = 3, lwd = 3)
```



# Heatmaps

```
set.seed(1234)
dataMatrix <- as.matrix(dataFrame)[sample(1:12), ]
kmeansObj2 <- kmeans(dataMatrix, centers = 3)
par(mfrow = c(1, 2), mar = c(2, 4, 0.1, 0.1))
image(t(dataMatrix)[, nrow(dataMatrix):1], yaxt = "n")
image(t(dataMatrix)[, order(kmeansObj2$cluster)], yaxt = "n")
```



# Notes and further resources

- K-means requires a number of clusters
  - Pick by eye/intuition
  - Pick by cross validation/information theory, etc.
  - [Determining the number of clusters](#)
- K-means is not deterministic
  - Different # of clusters
  - Different number of iterations
- [Rafael Irizarry's Distances and Clustering Video](#)
- [Elements of statistical learning](#)

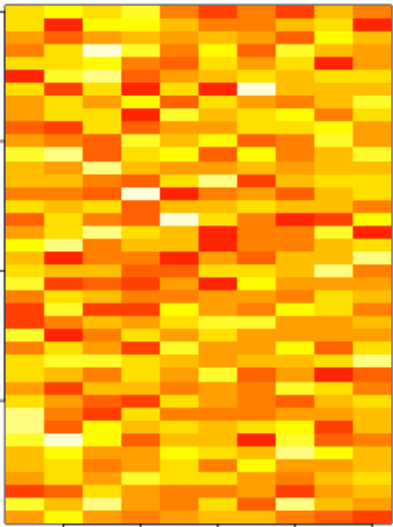


# Principal Components Analysis and Singular Value Decomposition

Roger D. Peng, Associate Professor of Biostatistics  
Johns Hopkins Bloomberg School of Public Health

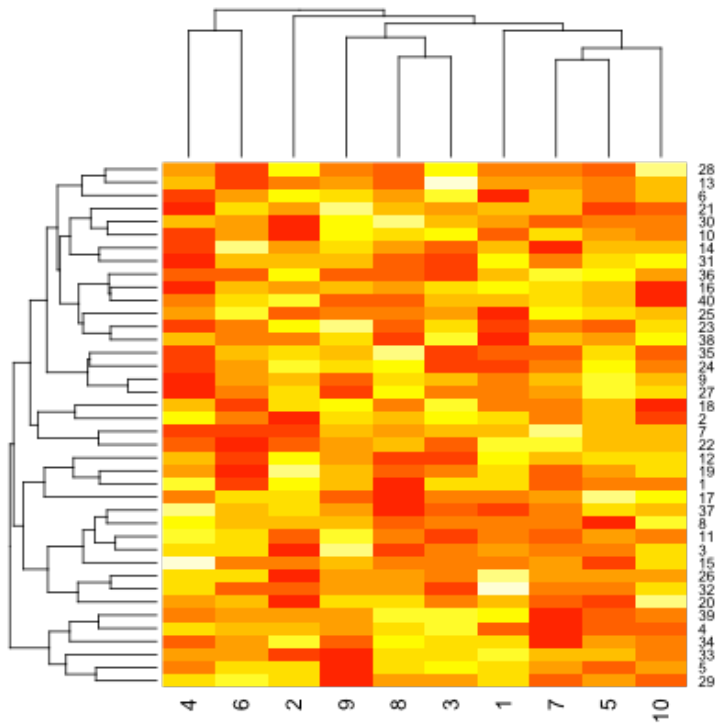
# Matrix data

```
set.seed(12345)
par(mar = rep(0.2, 4))
dataMatrix <- matrix(rnorm(400), nrow = 40)
image(1:10, 1:40, t(dataMatrix)[, nrow(dataMatrix):1])
```



# Cluster the data

```
par(mar = rep(0.2, 4))  
heatmap(dataMatrix)
```

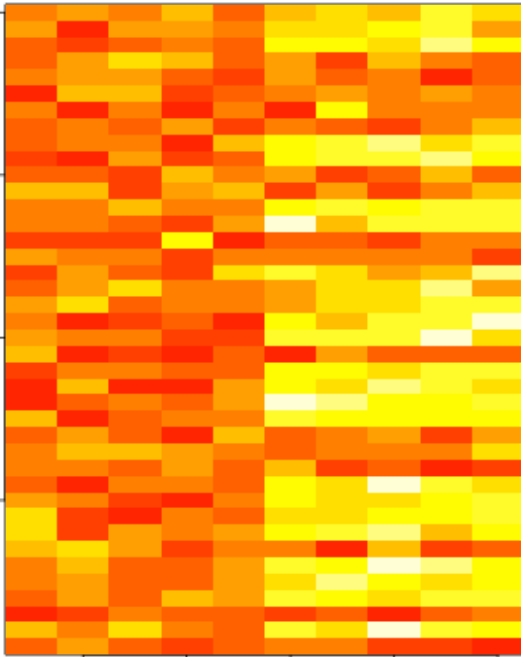


# What if we add a pattern?

```
set.seed(678910)
for (i in 1:40) {
  # flip a coin
  coinFlip <- rbinom(1, size = 1, prob = 0.5)
  # if coin is heads add a common pattern to that row
  if (coinFlip) {
    dataMatrix[i, ] <- dataMatrix[i, ] + rep(c(0, 3), each = 5)
  }
}
```

# What if we add a pattern? - the data

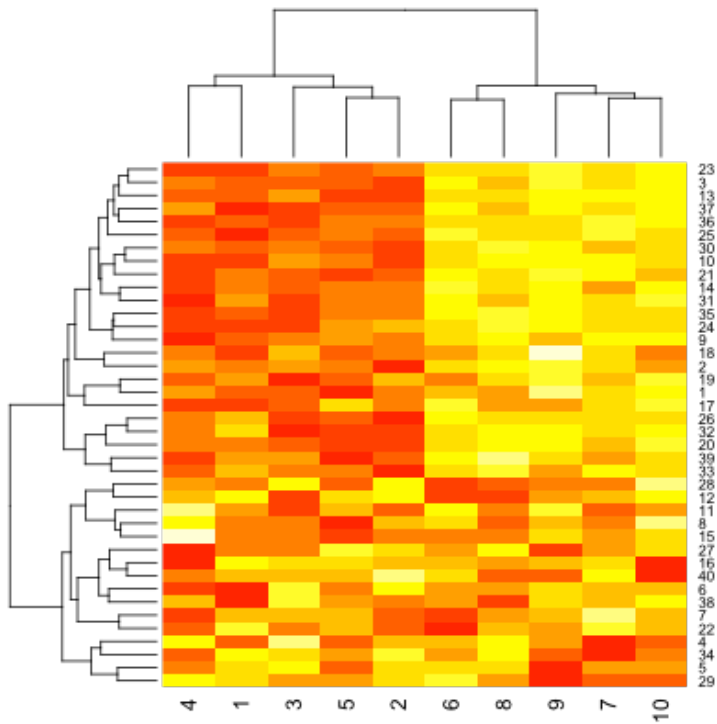
```
par(mar = rep(0.2, 4))  
image(1:10, 1:40, t(dataMatrix)[, nrow(dataMatrix):1])
```





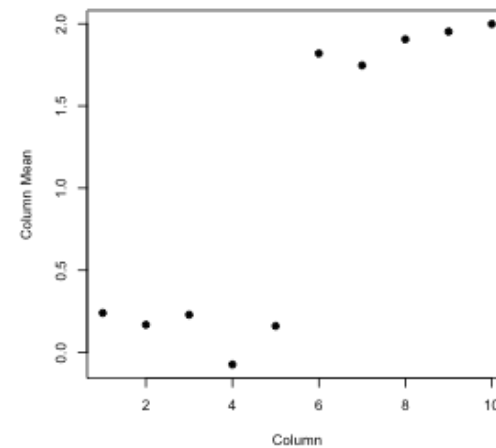
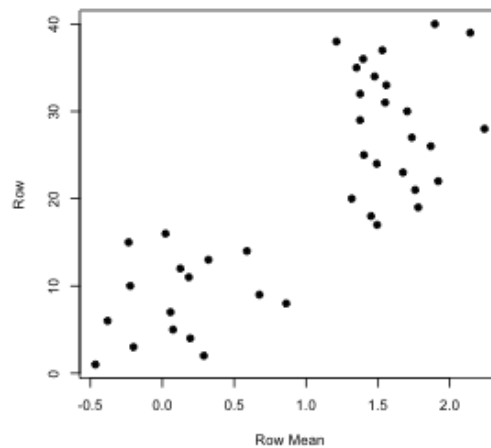
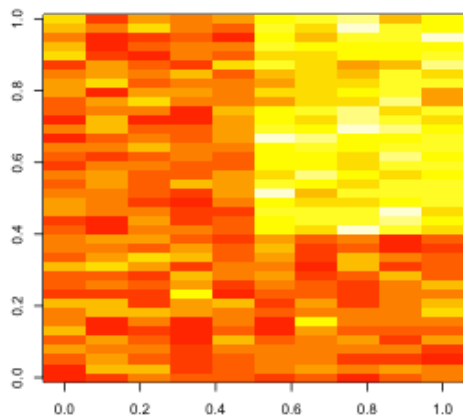
# What if we add a pattern? - the clustered data

```
par(mar = rep(0.2, 4))  
heatmap(dataMatrix)
```



# Patterns in rows and columns

```
hh <- hclust(dist(dataMatrix))
dataMatrixOrdered <- dataMatrix[hh$order, ]
par(mfrow = c(1, 3))
image(t(dataMatrixOrdered)[, nrow(dataMatrixOrdered):1])
plot(rowMeans(dataMatrixOrdered), 40:1, , xlab = "Row Mean", ylab = "Row", pch = 19)
plot(colMeans(dataMatrixOrdered), xlab = "Column", ylab = "Column Mean", pch = 19)
```



# Related problems

You have multivariate variables  $X_1, \dots, X_n$  so  $X_1 = (X_{11}, \dots, X_{1m})$

- Find a new set of multivariate variables that are uncorrelated and explain as much variance as possible.
- If you put all the variables together in one matrix, find the best matrix created with fewer variables (lower rank) that explains the original data.

The first goal is **statistical** and the second goal is **data compression**.

# Related solutions - PCA/SVD

## SVD

If  $X$  is a matrix with each variable in a column and each observation in a row then the SVD is a "matrix decomposition"

$$X = UDV^T$$

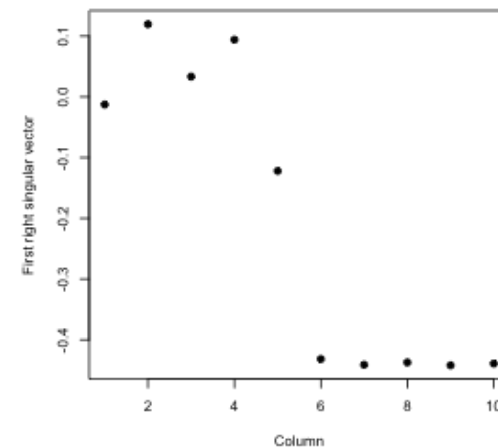
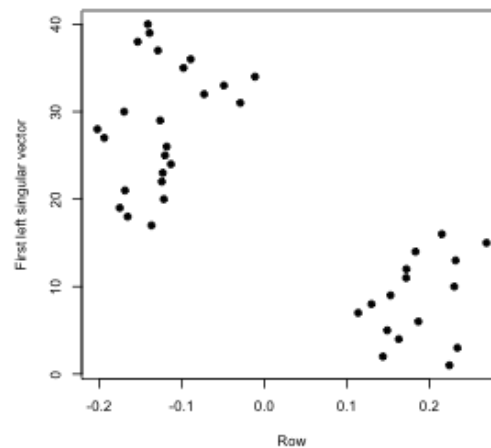
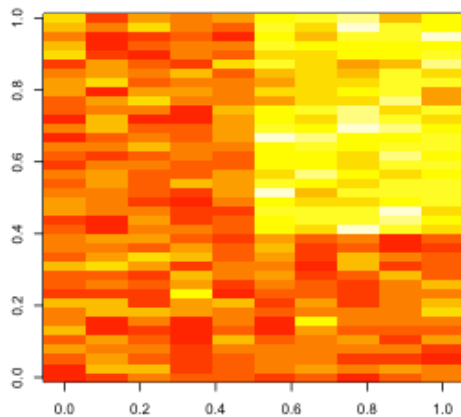
where the columns of  $U$  are orthogonal (left singular vectors), the columns of  $V$  are orthogonal (right singular vectors) and  $D$  is a diagonal matrix (singular values).

## PCA

The principal components are equal to the right singular values if you first scale (subtract the mean, divide by the standard deviation) the variables.

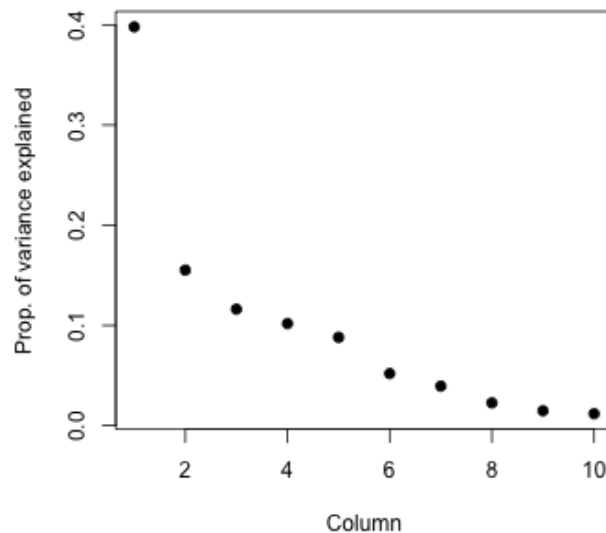
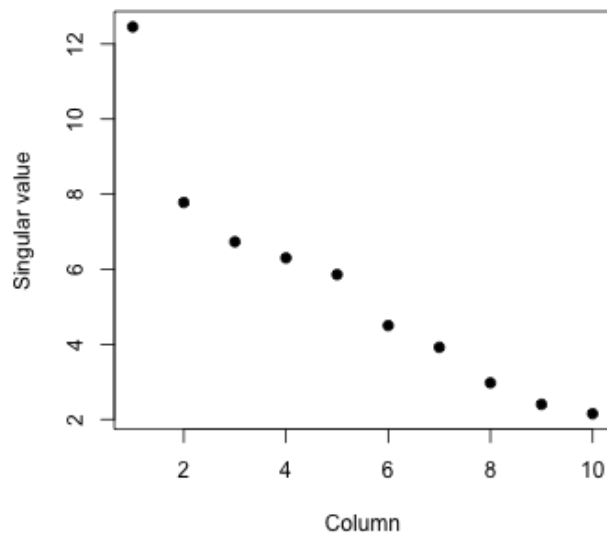
# Components of the SVD - $u$ and $v$

```
svd1 <- svd(scale(dataMatrixOrdered))  
par(mfrow = c(1, 3))  
image(t(dataMatrixOrdered)[, nrow(dataMatrixOrdered):1])  
plot(svd1$u[, 1], 40:1, , xlab = "Row", ylab = "First left singular vector",  
     pch = 19)  
plot(svd1$v[, 1], xlab = "Column", ylab = "First right singular vector", pch = 19)
```



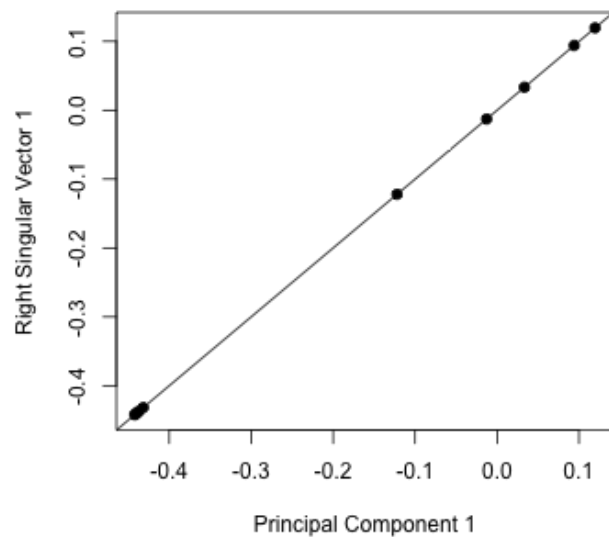
# Components of the SVD - Variance explained

```
par(mfrow = c(1, 2))  
plot(svd1$d, xlab = "Column", ylab = "Singular value", pch = 19)  
plot(svd1$d^2/sum(svd1$d^2), xlab = "Column", ylab = "Prop. of variance explained",  
     pch = 19)
```



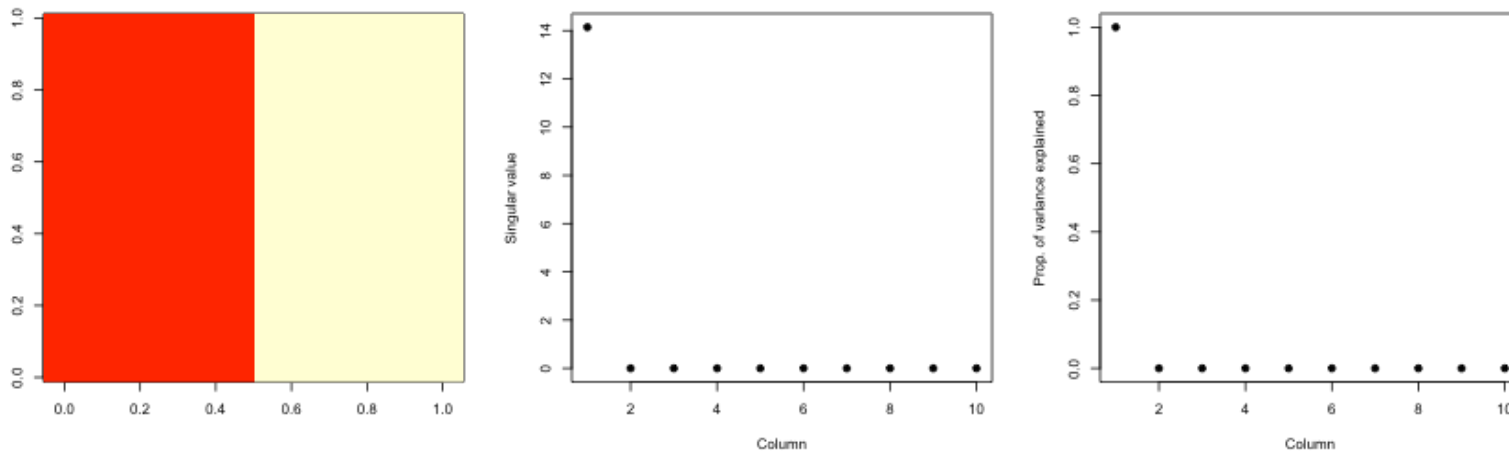
# Relationship to principal components

```
svd1 <- svd(scale(dataMatrixOrdered))  
pca1 <- prcomp(dataMatrixOrdered, scale = TRUE)  
plot(pca1$rotation[, 1], svd1$v[, 1], pch = 19, xlab = "Principal Component 1",  
      ylab = "Right Singular Vector 1")  
abline(c(0, 1))
```



# Components of the SVD - variance explained

```
constantMatrix <- dataMatrixOrdered*0
for(i in 1:dim(dataMatrixOrdered)[1]){constantMatrix[i,] <- rep(c(0,1),each=5)}
svd1 <- svd(constantMatrix)
par(mfrow=c(1,3))
image(t(constantMatrix)[,nrow(constantMatrix):1])
plot(svd1$d,xlab="Column",ylab="Singular value",pch=19)
plot(svd1$d^2/sum(svd1$d^2),xlab="Column",ylab="Prop. of variance explained",pch=19)
```



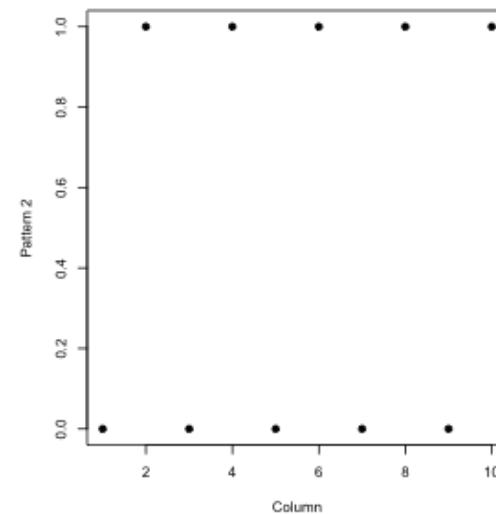
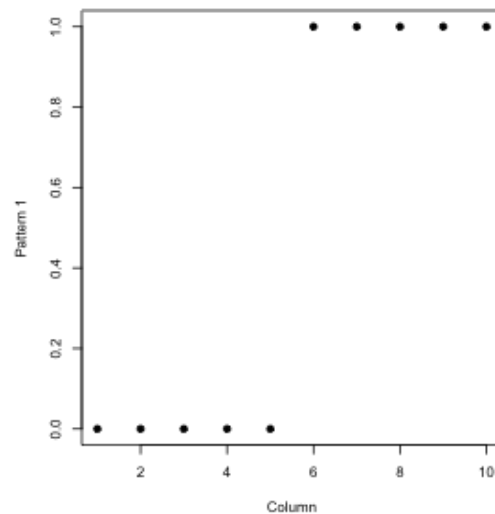
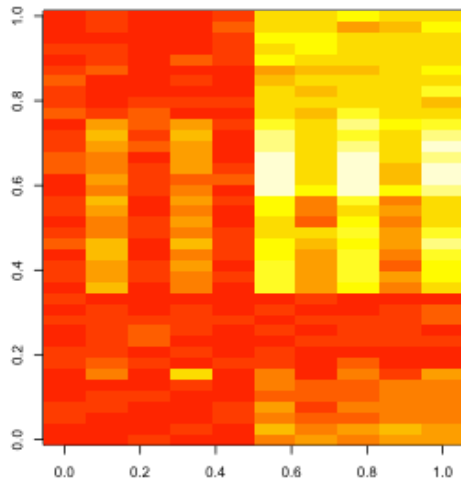


# What if we add a second pattern?

```
set.seed(678910)
for (i in 1:40) {
  # flip a coin
  coinFlip1 <- rbinom(1, size = 1, prob = 0.5)
  coinFlip2 <- rbinom(1, size = 1, prob = 0.5)
  # if coin is heads add a common pattern to that row
  if (coinFlip1) {
    dataMatrix[i, ] <- dataMatrix[i, ] + rep(c(0, 5), each = 5)
  }
  if (coinFlip2) {
    dataMatrix[i, ] <- dataMatrix[i, ] + rep(c(0, 5), 5)
  }
}
hh <- hclust(dist(dataMatrix))
dataMatrixOrdered <- dataMatrix[hh$order, ]
```

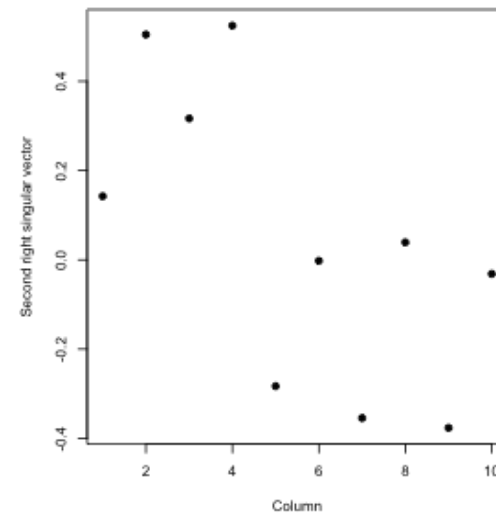
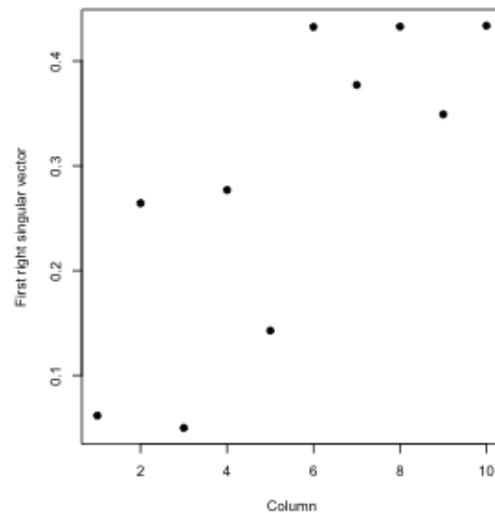
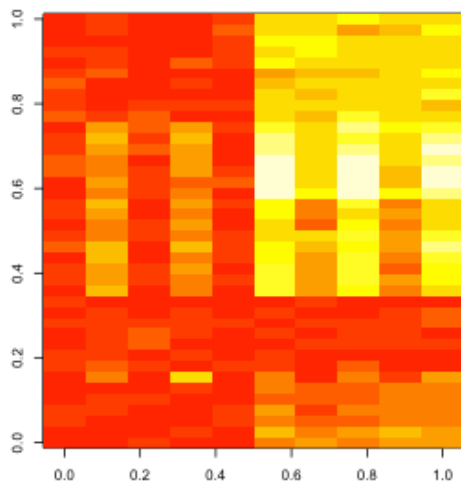
# Singular value decomposition - true patterns

```
svd2 <- svd(scale(dataMatrixOrdered))  
par(mfrow = c(1, 3))  
image(t(dataMatrixOrdered)[, nrow(dataMatrixOrdered):1])  
plot(rep(c(0, 1), each = 5), pch = 19, xlab = "Column", ylab = "Pattern 1")  
plot(rep(c(0, 1), 5), pch = 19, xlab = "Column", ylab = "Pattern 2")
```



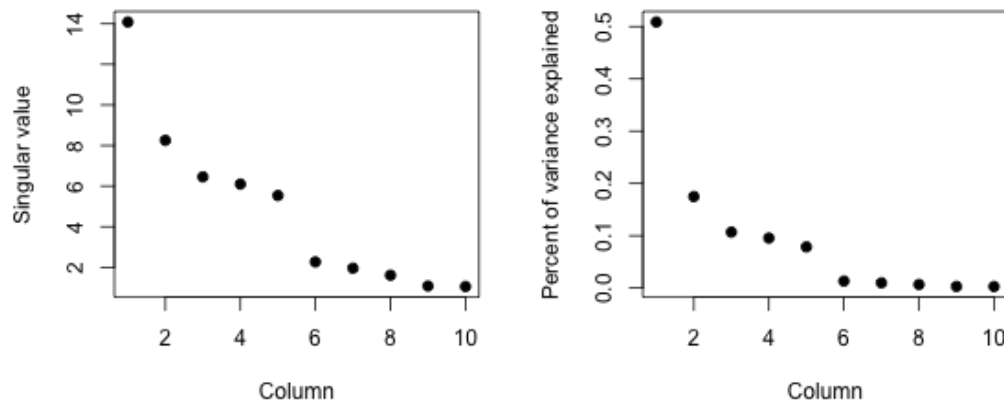
# $v$ and patterns of variance in rows

```
svd2 <- svd(scale(dataMatrixOrdered))  
par(mfrow = c(1, 3))  
image(t(dataMatrixOrdered)[, nrow(dataMatrixOrdered):1])  
plot(svd2$v[, 1], pch = 19, xlab = "Column", ylab = "First right singular vector")  
plot(svd2$v[, 2], pch = 19, xlab = "Column", ylab = "Second right singular vector")
```



# $d$ and variance explained

```
svd1 <- svd(scale(dataMatrixOrdered))  
par(mfrow = c(1, 2))  
plot(svd1$d, xlab = "Column", ylab = "Singular value", pch = 19)  
plot(svd1$d^2/sum(svd1$d^2), xlab = "Column", ylab = "Percent of variance explained",  
     pch = 19)
```



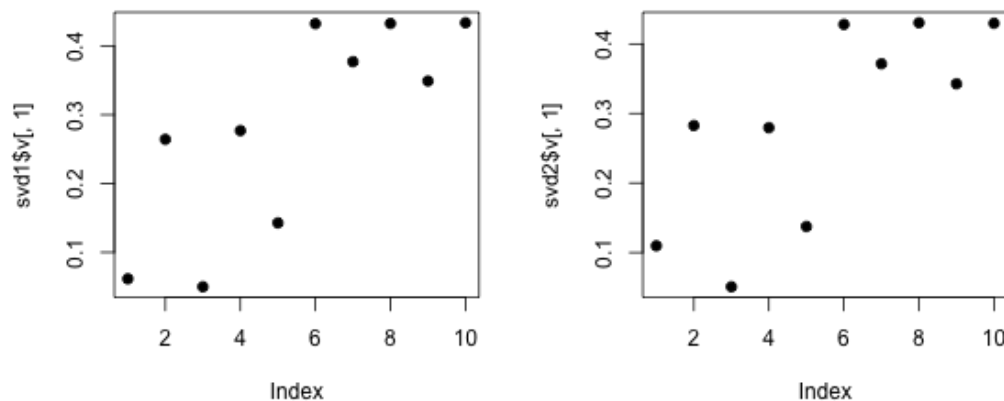
# Missing values

```
dataMatrix2 <- dataMatrixOrdered  
## Randomly insert some missing data  
dataMatrix2[sample(1:100, size = 40, replace = FALSE)] <- NA  
svd1 <- svd(scale(dataMatrix2)) ## Doesn't work!
```

```
## Error: infinite or missing values in 'x'
```

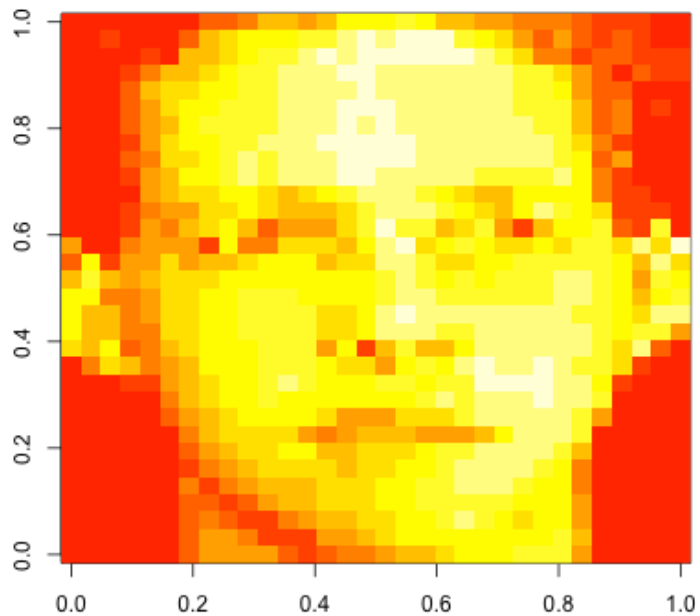
# Imputing {impute}

```
library(impute) ## Available from http://bioconductor.org
dataMatrix2 <- dataMatrixOrdered
dataMatrix2[sample(1:100,size=40,replace=FALSE)] <- NA
dataMatrix2 <- impute.knn(dataMatrix2)$data
svd1 <- svd(scale(dataMatrixOrdered)); svd2 <- svd(scale(dataMatrix2))
par(mfrow=c(1,2)); plot(svd1$v[,1],pch=19); plot(svd2$v[,1],pch=19)
```



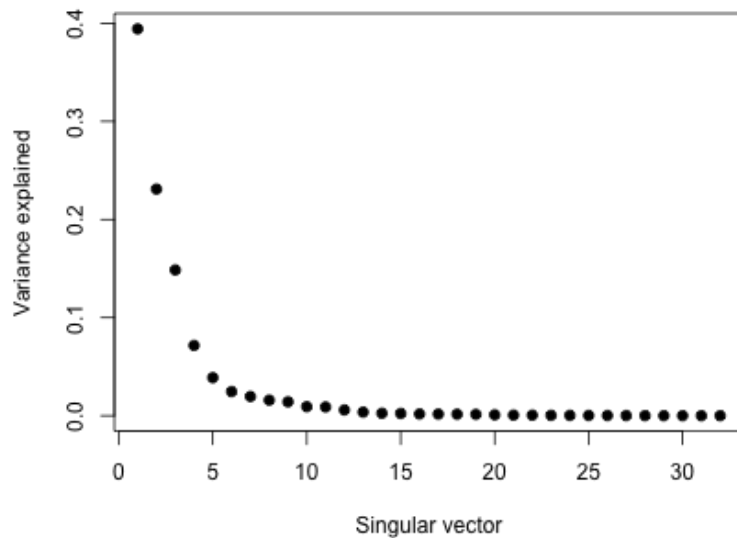
# Face example

```
load("data/face.rda")  
image(t(faceData)[, nrow(faceData):1])
```



# Face example - variance explained

```
svd1 <- svd(scale(faceData))  
plot(svd1$d^2/sum(svd1$d^2), pch = 19, xlab = "Singular vector", ylab = "Variance explained")
```



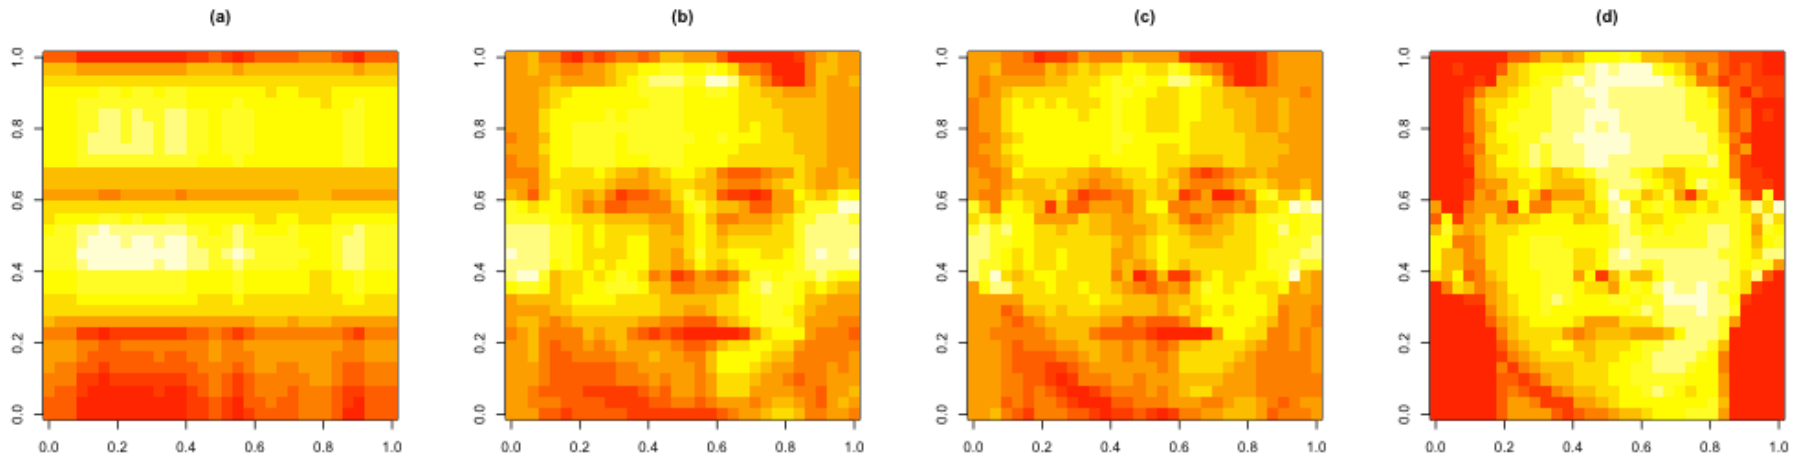


# Face example - create approximations

```
svd1 <- svd(scale(faceData))  
## Note that %*% is matrix multiplication  
  
# Here svd1$d[1] is a constant  
approx1 <- svd1$u[, 1] %*% t(svd1$v[, 1]) * svd1$d[1]  
  
# In these examples we need to make the diagonal matrix out of d  
approx5 <- svd1$u[, 1:5] %*% diag(svd1$d[1:5]) %*% t(svd1$v[, 1:5])  
approx10 <- svd1$u[, 1:10] %*% diag(svd1$d[1:10]) %*% t(svd1$v[, 1:10])
```

# Face example - plot approximations

```
par(mfrow = c(1, 4))  
image(t(approx1)[, nrow(approx1):1], main = "(a)")  
image(t(approx5)[, nrow(approx5):1], main = "(b)")  
image(t(approx10)[, nrow(approx10):1], main = "(c)")  
image(t(faceData)[, nrow(faceData):1], main = "(d)") ## Original data
```



# Notes and further resources

- Scale matters
- PC's/SV's may mix real patterns
- Can be computationally intensive
- [Advanced data analysis from an elementary point of view](#)
- [Elements of statistical learning](#)
- Alternatives
  - [Factor analysis](#)
  - [Independent components analysis](#)
  - [Latent semantic analysis](#)



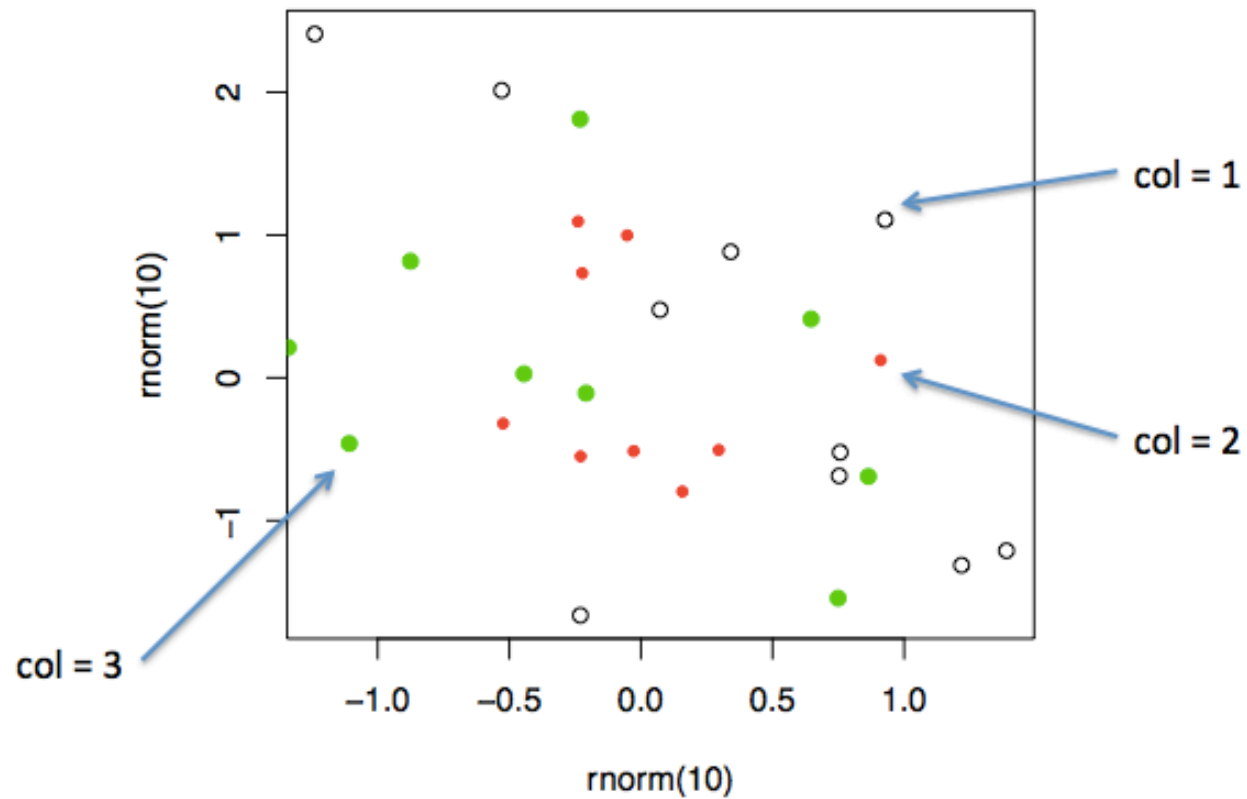
# Plotting and Color in R

Roger D. Peng, Associate Professor of Biostatistics  
Johns Hopkins Bloomberg School of Public Health

# Plotting and Color

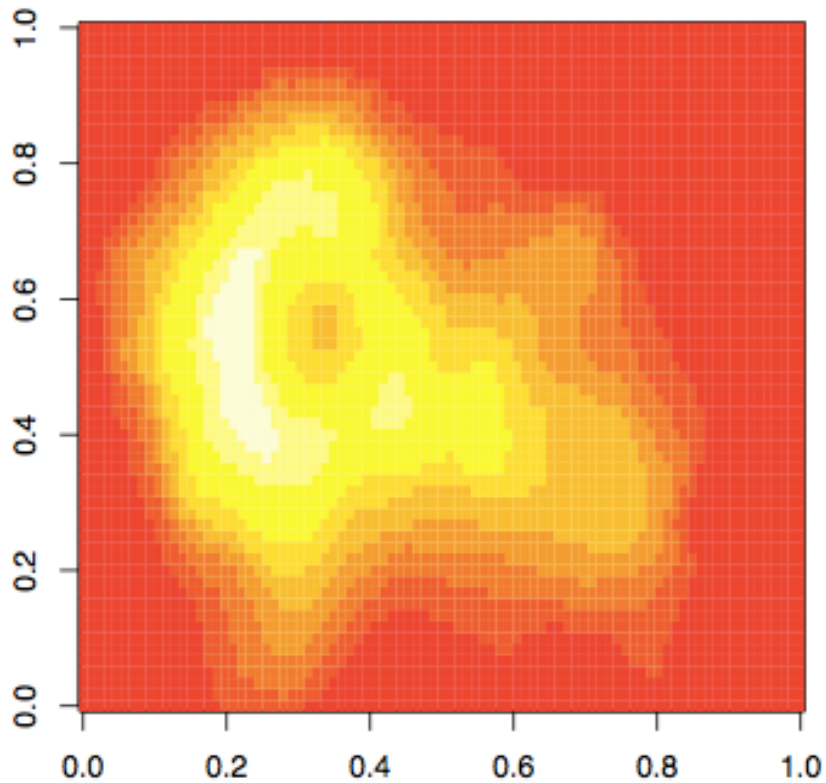
- The default color schemes for most plots in R are horrendous
  - I don't have good taste and even I know that
- Recently there have been developments to improve the handling/specification of colors in plots/graphs/etc.
- There are functions in R and in external packages that are very handy

# Colors 1, 2, and 3

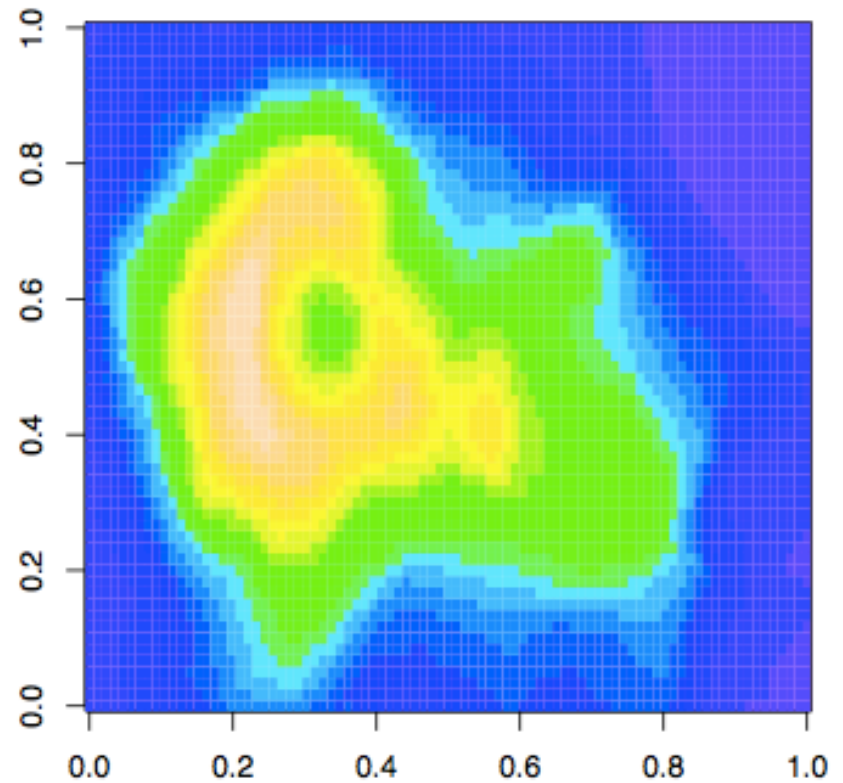


# Default Image Plots in R

heat.colors()



topo.colors()



# Color Utilities in R

- The `grDevices` package has two functions
  - `colorRamp`
  - `colorRampPalette`
- These functions take palettes of colors and help to interpolate between the colors
- The function `colors()` lists the names of colors you can use in any plotting function



# Color Palette Utilities in R

- `colorRamp`: Take a palette of colors and return a function that takes values between 0 and 1, indicating the extremes of the color palette (e.g. see the 'gray' function)
- `colorRampPalette`: Take a palette of colors and return a function that takes integer arguments and returns a vector of colors interpolating the palette (like `heat.colors` or `topo.colors`)

# colorRamp

[,1] [,2] [,3] corresponds to [Red] [Blue] [Green]

```
> pal <- colorRamp(c("red", "blue"))
```

```
> pal(0)
```

```
      [,1] [,2] [,3]  
[1,]  255    0    0
```

```
> pal(1)
```

```
      [,1] [,2] [,3]  
[1,]    0    0  255
```

```
> pal(0.5)
```

```
      [,1] [,2] [,3]  
[1,] 127.5    0 127.5
```

# colorRamp

```
> pal(seq(0, 1, len = 10))
```

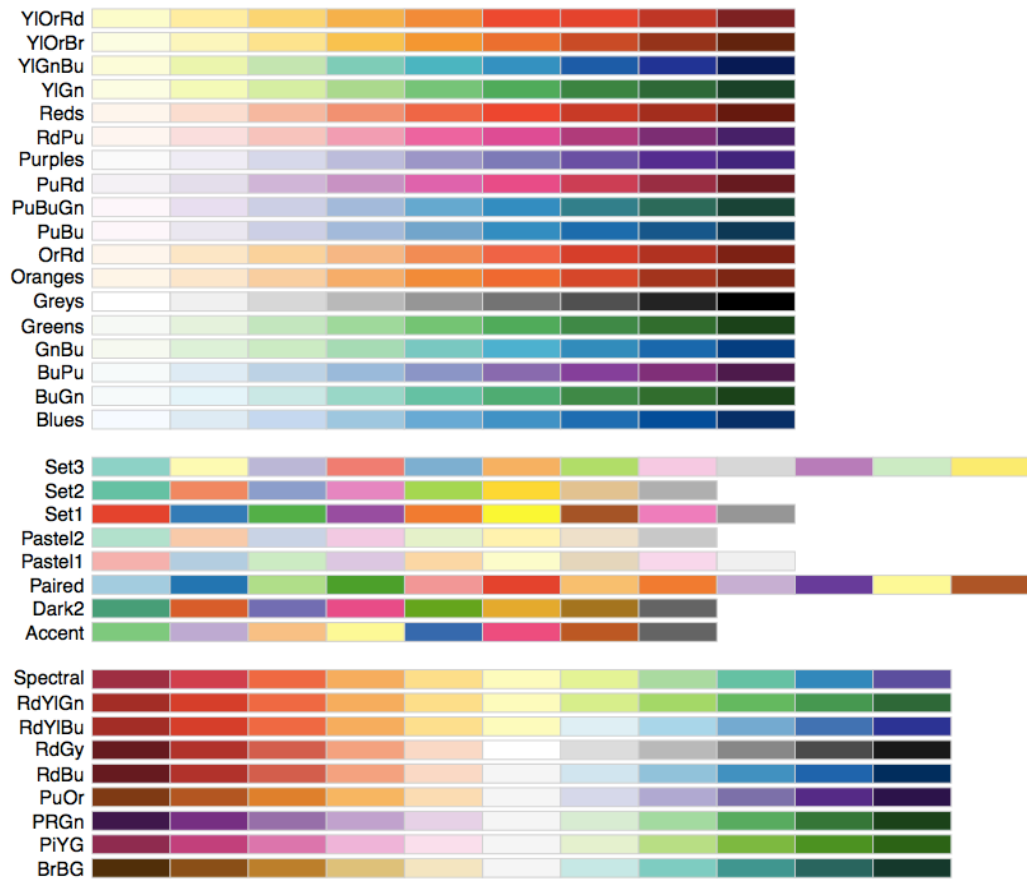
	[,1]	[,2]	[,3]
[1,]	255.00000	0	0
[2,]	226.66667	0	28.33333
[3,]	198.33333	0	56.66667
[4,]	170.00000	0	85.00000
[5,]	141.66667	0	113.33333
[6,]	113.33333	0	141.66667
[7,]	85.00000	0	170.00000
[8,]	56.66667	0	198.33333
[9,]	28.33333	0	226.66667
[10,]	0.00000	0	255.00000

# colorRampPalette

```
> pal <- colorRampPalette(c("red", "yellow"))  
  
> pal(2)  
[1] "#FF0000" "#FFFF00"  
  
> pal(10)  
[1] "#FF0000" "#FF1C00" "#FF3800" "#FF5500" "#FF7100"  
[6] "#FF8D00" "#FFAA00" "#FFC600" "#FFE200" "#FFFF00"
```

# RColorBrewer Package

- One package on CRAN that contains interesting/useful color palettes
- There are 3 types of palettes
  - Sequential
  - Diverging
  - Qualitative
- Palette information can be used in conjunction with the `colorRamp()` and `colorRampPalette()`



# RColorBrewer and colorRampPalette

```
> library(RColorBrewer)

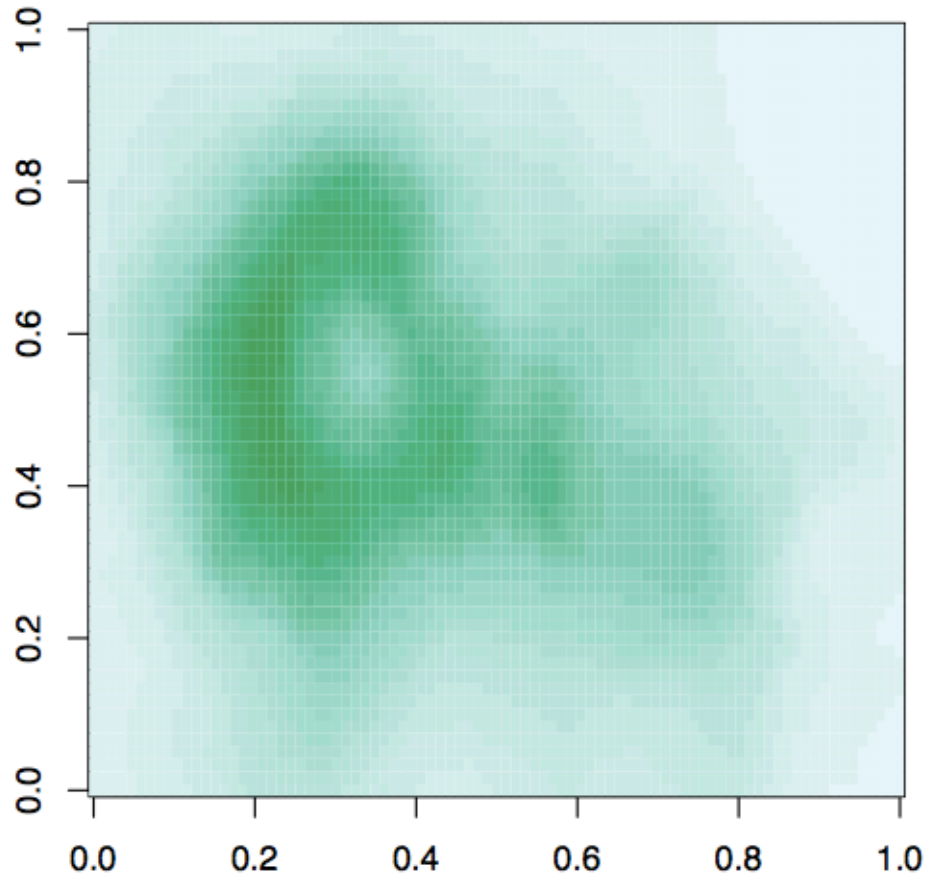
> cols <- brewer.pal(3, "BuGn")

> cols
[1] "#E5F5F9" "#99D8C9" "#2CA25F"

> pal <- colorRampPalette(cols)

> image(volcano, col = pal(20))
```

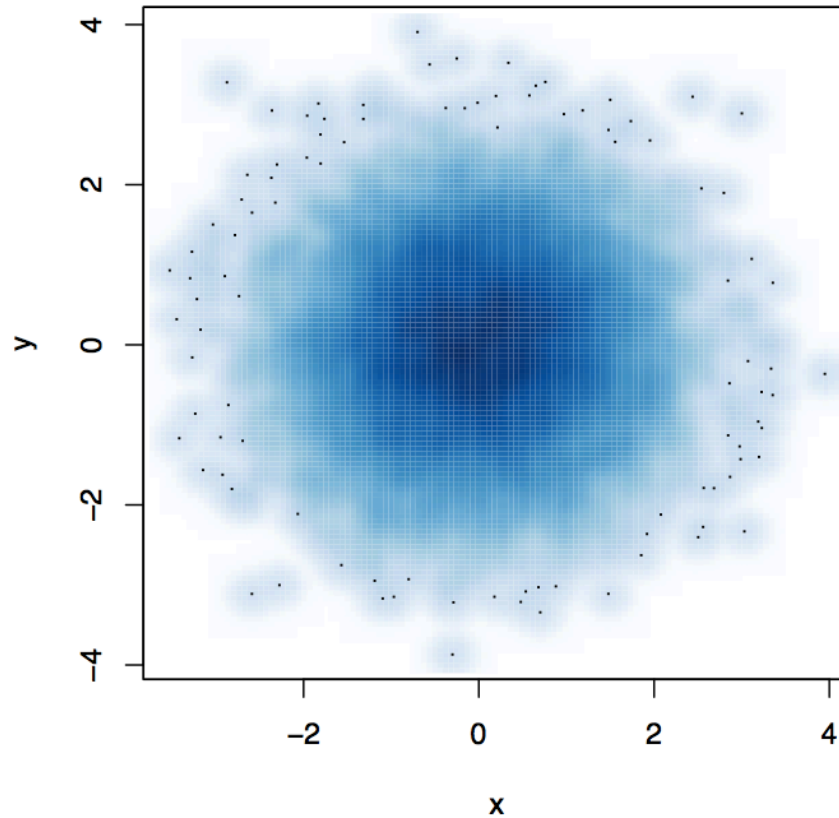
# RColorBrewer and colorRampPalette





# The smoothScatter function

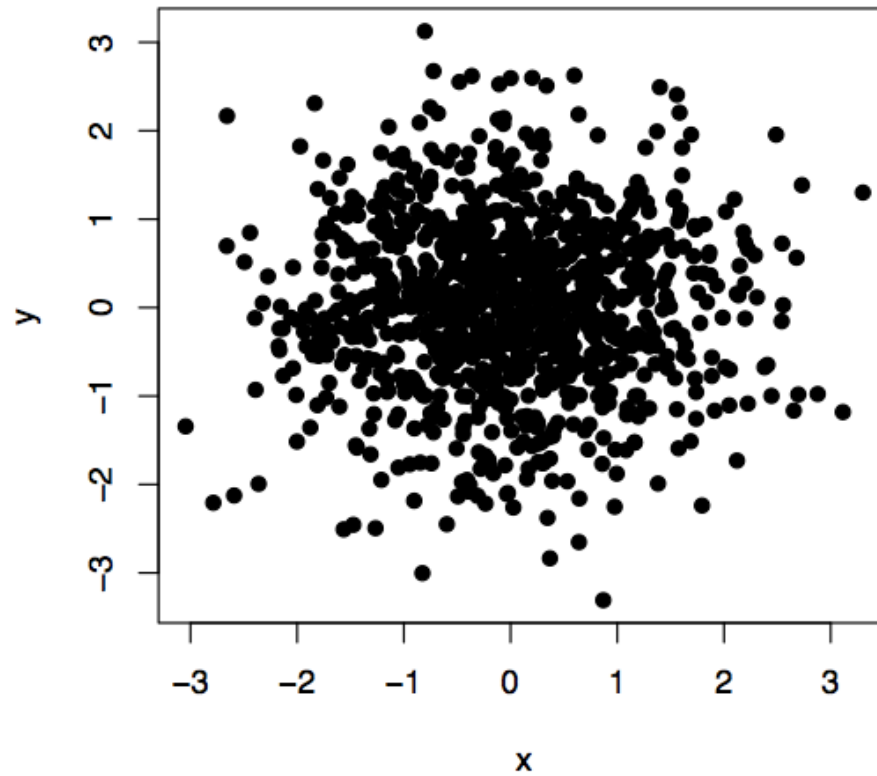
```
x <- rnorm(10000)  
y <- rnorm(10000)  
smoothScatter(x, y)
```



# Some other plotting notes

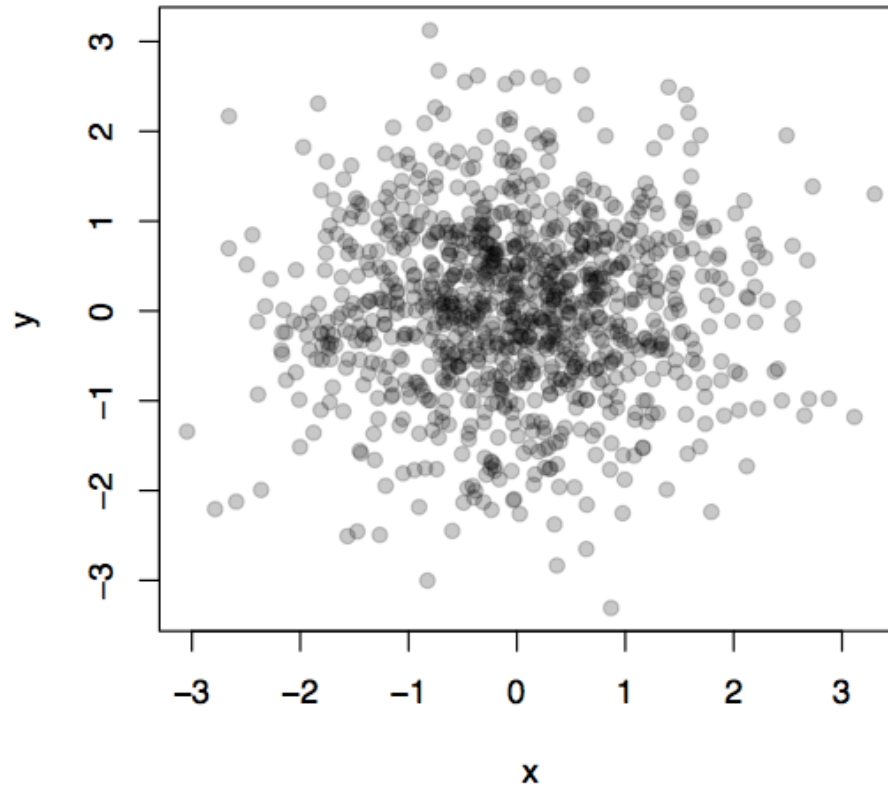
- The `rgb` function can be used to produce any color via red, green, blue proportions
- Color transparency can be added via the `alpha` parameter to `rgb`
- The `colorspace` package can be used for a different control over colors

# Scatterplot with no transparency



```
plot(x, y, pch = 19)
```

# Scatterplot with transparency



```
plot(x, y, col = rgb(0, 0, 0, 0.2), pch = 19)
```

# Summary

- Careful use of colors in plots/maps/etc. can make it easier for the reader to get what you're trying to say (why make it harder?)
- The `RColorBrewer` package is an R package that provides color palettes for sequential, categorical, and diverging data
- The `colorRamp` and `colorRampPalette` functions can be used in conjunction with color palettes to connect data to colors
- Transparency can sometimes be used to clarify plots with many points